

1	a) Create 2 text files. Read the contents in a single RDD. b) Create 2 CSV files. Read the contents in a single RDD.	20
2	Create two dataframes one for employee and other for dept. Perform a) Left outer join b) Full outer join c) Inner join	20
4	Viva	5
5	Journal	5

1

Step 1: Install PySpark and set up the Spark Session

```
!pip install pyspark
```

```
from pyspark.sql import SparkSession
```

```
spark = SparkSession.builder.appName("PracticalExam_Q1").getOrCreate()
```

```
sc = spark.sparkContext
```

```
print("--- Spark Session Created ---\n")
```

Step 2: Create the necessary files for the question

```
with open("file1.txt", "w") as f:
```

```
    f.write("this is the first text file\n")
```

```
    f.write("it has two lines\n")
```

```
with open("file2.txt", "w") as f:
```

```
    f.write("this is the second text file\n")
```

```
with open("data1.csv", "w") as f:
```

```
    f.write("id,value\n")
```

```
    f.write("1,a\n")
```

```
    f.write("2,b\n")
```

```
with open("data2.csv", "w") as f:
```

```
    f.write("id,value\n")
```

```
    f.write("3,c\n")
```

```
    f.write("4,d\n")
```

```
print("--- Sample files for Question 1 created successfully ---\n")
```

--- SOLUTION ---

a) Read 2 text files into a single RDD

```
print("\nReading contents from two text files into a single RDD:")
```

```
text_rdd = sc.textFile("file1.txt,file2.txt")
```

```
print("Result:", text_rdd.collect())
```

a) Read 2 CSV files into a single RDD

```
print("\nReading contents from two CSV files into a single RDD:")
```

```
csv_df = spark.read.csv(["data1.csv", "data2.csv"], header=True)
csv_rdd = csv_df.rdd
print("Result:", csv_rdd.collect())
```

2.

Assuming the SparkSession 'spark' is already created from the previous question.

--- SETUP ---

Create the employee and department DataFrames

```
emp_data = [(1, "Smith", 10), (2, "Rose", 20), (3, "Williams", 10), (4, "Jones", 30)]
```

```
dept_data = [("Finance", 10), ("Marketing", 20), ("Sales", 30), ("IT", 40)]
```

```
emp_df = spark.createDataFrame(emp_data, ["emp_id", "name", "dept_id"])
```

```
dept_df = spark.createDataFrame(dept_data, ["dept_name", "dept_id"])
```

```
print("Employee DataFrame:")
```

```
emp_df.show()
```

```
print("Department DataFrame:")
```

```
dept_df.show()
```

--- SOLUTION ---

a) Perform Left outer join

```
print("\na) Left Outer Join Result:")
```

```
emp_df.join(dept_df, on="dept_id", how="left_outer").show()
```

b) Perform Full outer join

```
print("\nb) Full Outer Join Result:")
```

```
emp_df.join(dept_df, on="dept_id", how="full_outer").show()
```

c) Perform Inner join

```
print("\nc) Inner Join Result:")
```

```
emp_df.join(dept_df, on="dept_id", how="inner").show()
```

1	<p>For the following data and schema create a dataframe and perform the given operations</p> <p>Data: Seq(Row(Row("James;", "", "Smith"), "36636", "M", "20000"), Row(Row("Michael", "Rose", ""), "40288", "M", "40000"), Row(Row("Robert", "", "Williams"), "42114", "M", "10000"), Row(Row("Maria", "Anne", "Jones"), "39192", "F", "45000"), Row(Row("Jen", "Mary", "Brown"), "", "F", "-1"))</p> <p>Schema should have the columns as: firstname, middlename, lastname, dob, gender, expenses All columns will be of type String</p> <p>Perform the following operations:</p> <ol style="list-style-type: none"> Change the data type of expenses to Integer Rename dob to DateOfBirth Create a column that has value expense*5 	20
2	<p>Create a data frame with a nested array column. Perform the following operations:</p> <ol style="list-style-type: none"> Flatten nested array Explode nested array Convert array of string to string column. 	20
3	Viva	5
4	Journal	5

1.

Step 1: Install PySpark and set up the Spark Session

```
!pip install pyspark
```

```
from pyspark.sql import SparkSession, Row
```

```
from pyspark.sql.types import StructType, StructField, StringType, IntegerType
```

```
from pyspark.sql.functions import col
```

```
spark = SparkSession.builder.appName("PracticalExam_Slip2_Q1").getOrCreate()
```

```
print("--- Spark Session Created ---\n")
```

```
# --- SETUP ---
```

```
# Define the data and schema
```

```
data = [
```

```
    Row(name=Row(firstname="James", middlename="", lastname="Smith"), dob="36636", gender="M",  
expenses="20000"),
```

```
    Row(name=Row(firstname="Michael", middlename="Rose", lastname=""), dob="40288", gender="M",  
expenses="40000"),
```

```
    Row(name=Row(firstname="Robert", middlename="", lastname="Williams"), dob="42114", gender="M",  
expenses="10000"),
```

```
    Row(name=Row(firstname="Maria", middlename="Anne", lastname="Jones"), dob="39192", gender="F",  
expenses="45000"),
```

```
    Row(name=Row(firstname="Jen", middlename="Mary", lastname="Brown"), dob="", gender="F", expenses="-  
1")
```

```
]
```

```

schema = StructType([
    StructField("name", StructType([
        StructField("firstname", StringType(), True),
        StructField("middlename", StringType(), True),
        StructField("lastname", StringType(), True)
    ])),
    StructField("dob", StringType(), True),
    StructField("gender", StringType(), True),
    StructField("expenses", StringType(), True)
])

```

```

# Create the DataFrame
df = spark.createDataFrame(data, schema)
print("Original DataFrame:")
df.show(truncate=False)
df.printSchema()

```

--- SOLUTION ---

```

# a) Change the data type of expenses to Integer
print("\na) Changing 'expenses' to Integer type:")
df_a = df.withColumn("expenses", col("expenses").cast(IntegerType()))
df_a.printSchema()
df_a.show()

```

```

# b) Rename dob to DateOfBirth
print("\nb) Renaming 'dob' to 'DateOfBirth':")
df_b = df_a.withColumnRenamed("dob", "DateOfBirth")
df_b.show()

```

```

# c) Create a column that has value expense*5
print("\nc) Creating a 'bonus' column with 'expenses * 5':")
df_c = df_b.withColumn("bonus", col("expenses") * 5)
df_c.show()

```

2.

```

# Assuming the SparkSession 'spark' is already created from the previous question.
from pyspark.sql.functions import col, flatten, explode, concat_ws
from pyspark.sql.types import StructType, StructField, StringType, ArrayType

```

--- SETUP ---

```

# Create the DataFrame with a nested array
data = [
    ("James", [ ["Java", "Scala", "C++"], ["Spark", "Java"] ]),
    ("Michael", [ ["Spark", "Java", "C++"], ["Spark", "Java"] ]),
    ("Robert", [ ["CSharp", "VB"], ["Spark", "Python"] ])
]

schema = StructType([
    StructField("name", StringType(), True),
    StructField("subjects", ArrayType(ArrayType(StringType()), True)

```

```
])  
df = spark.createDataFrame(data, schema)  
print("Original DataFrame with nested array:")  
df.show(truncate=False)
```

--- SOLUTION ---

```
# a) Flatten nested array  
print("\na) Flattened nested array:")  
df_a = df.withColumn("subjects_flat", flatten(col("subjects")))  
df_a.show(truncate=False)
```

```
# b) Explode nested array  
print("\nb) Exploded nested array:")  
# Note: Exploding a nested array directly creates rows with the inner arrays.  
# To explode to individual elements, you must flatten first.  
df_b = df_a.withColumn("subject", explode(col("subjects_flat")))  
df_b.show(truncate=False)
```

```
# c) Convert array of string to string column  
print("\nc) Converted array to a single string column:")  
df_c = df_a.withColumn("subjects_string", concat_ws(" ", col("subjects_flat")))  
df_c.show(truncate=False)
```

1	a) Create a data frame with today's date and timestamp b) Display the hours, minutes and seconds from the timestamp	20
2	For the following employee data showing name, dept and salary, perform the given operations: Data: ("James", "Sales", 3000), ("Michael", "Sales", 4600), ("Robert", "Sales", 4100), ("Maria", "Finance", 3000), ("James", "Sales", 3000), ("Scott", "Finance", 3300), ("Jen", "Finance", 3900), ("Jeff", "Marketing", 3000), ("Kumar", "Marketing", 2000), ("Saif", "Sales", 4100), ("Jason", "Sales", 9000), ("Alice", "Finance", 3700), ("Jenniffer", "Finance", 8900), ("Jenson", "Marketing", 9000) a) Create a data frame for the above data b) Display average salary c) Display number of unique departments d) Display number of employees with unique salary	20
3	Viva	5
4	Journal	5

1

Step 1: Install PySpark and set up the Spark Session

```
!pip install pyspark
```

```
from pyspark.sql import SparkSession
```

```
from pyspark.sql.functions import current_date, current_timestamp, hour, minute, second, col
```

```
spark = SparkSession.builder.appName("PracticalExam_Slip3_Q1").getOrCreate()
```

```
print("--- Spark Session Created ---\n")
```

--- SOLUTION ---

a) Create a data frame with today's date and timestamp

We start with a dummy DataFrame with one row to add columns to.

```
df = spark.range(1)
```

```
df_with_time = df.withColumn("today_date", current_date()) \
    .withColumn("current_ts", current_timestamp())
```

```
print("a) DataFrame with current date and timestamp:")
```

```
df_with_time.show(truncate=False)
```

b) Display the hours, minutes and seconds from the timestamp

```

time_parts_df = df_with_time.withColumn("hour", hour(col("current_ts"))) \
    .withColumn("minute", minute(col("current_ts"))) \
    .withColumn("second", second(col("current_ts")))

print("\nb) Timestamp parts extracted:")
time_parts_df.select("current_ts", "hour", "minute", "second").show(truncate=False)

```

2.

```

# Assuming the SparkSession 'spark' is already created from the previous question.
from pyspark.sql.functions import avg, countDistinct

```

```

# --- SETUP ---

```

```

# The provided data has a typo `('Jason'`, which has been corrected to `('Jason'`.

```

```

employee_data = [
    ("James", "Sales", 3000), ("Michael", "Sales", 4600),
    ("Robert", "Sales", 4100), ("Maria", "Finance", 3000),
    ("James", "Sales", 3000), ("Scott", "Finance", 3300),
    ("Jen", "Finance", 3900), ("Jeff", "Marketing", 3000),
    ("Kumar", "Marketing", 2000), ("Saif", "Sales", 4100),
    ("Jason", "Sales", 9000), ("Alice", "Finance", 3700),
    ("Jenniffer", "Finance", 8900), ("Jenson", "Marketing", 9000)
]

columns = ["name", "department", "salary"]

```

```

# --- SOLUTION ---

```

```

# a) Create a data frame for the above data

```

```

emp_df = spark.createDataFrame(employee_data, columns)
print("\na) Employee DataFrame:")
emp_df.show()

```

```

# b) Display average salary

```

```

avg_salary_df = emp_df.select(avg("salary").alias("average_salary"))
print("\nb) Average salary:")
avg_salary_df.show()

```

```

# c) Display number of unique departments

```

```

unique_dept_df = emp_df.select(countDistinct("department").alias("unique_departments"))
print("\nc) Number of unique departments:")
unique_dept_df.show()

```

```

# d) Display number of employees with unique salary

```

```

# This is interpreted as the count of distinct salary values.

```

```

unique_salary_df = emp_df.select(countDistinct("salary").alias("unique_salary_count"))
print("\nd) Number of unique salary values:")
unique_salary_df.show()

```

1	a) Create a data frame containing today's date, date 2022-01-31, date 2021-03-22, date 2024-01-31, date 2023-11-11. b) Store the date in the format MM-DD-YYYY. c) Display the dates in the format DD/MM/YYYY d) Find the number of months between each of the dates and today's date	20
2	a) Create data frame with a column that contains JSON string. b) Convert the JSON string into Struct type or Map type. c) Extract the Data from JSON and create them as new columns. d) Convert MapType or Struct type to JSON string	20
3	Viva	5
4	Journal	5

1.

Step 1: Install PySpark and set up the Spark Session

```
!pip install pyspark
```

```
from pyspark.sql import SparkSession
```

```
from pyspark.sql.functions import col, current_date, to_date, date_format, lit, months_between
```

```
from pyspark.sql.types import StructType, StructField, StringType
```

```
spark = SparkSession.builder.appName("PracticalExam_Slip4_Q1").getOrCreate()
```

```
print("--- Spark Session Created ---\n")
```

--- SOLUTION ---

a) and b) Create a DataFrame with dates stored in MM-DD-YYYY format

```
date_data = [
```

```
    ("01-31-2022",),
```

```
    ("03-22-2021",),
```

```
    ("01-31-2024",),
```

```
    ("11-11-2023",)
```

```
]
```

Add today's date to the list

```
today's_date_str = spark.range(1).select(date_format(current_date(), "MM-dd-yyyy")).first()[0]
```

```
date_data.append((today's_date_str,))
```

```
date_df = spark.createDataFrame(date_data, ["date_str_mmddyyyy"])
```

```
print("a) and b) DataFrame with dates in MM-DD-YYYY format:")
```

```
date_df.show()
```

To perform date operations, first convert the strings to a proper DateType

```
df_with_dates = date_df.withColumn("date_obj", to_date(col("date_str_mmddyyyy"), "MM-dd-yyyy"))
```

c) Display the dates in the format DD/MM/YYYY

```
print("\nc) Dates displayed in DD/MM/YYYY format:")
```

```
df_formatted = df_with_dates.withColumn("date_str_ddmmyyyy", date_format(col("date_obj"), "dd/MM/yyyy"))
```

```
df_formatted.select("date_str_mmddyyyy", "date_str_ddmmyyyy").show()
```

d) Find the number of months between each date and today's date

```
print("\nd) Number of months between each date and today:")
```

```
df_months_between = df_with_dates.withColumn("months_from_today", months_between(current_date(), col("date_obj")))
```

```
df_months_between.select("date_str_mmddyyyy", "months_from_today").show()
```


2.

Assuming the SparkSession 'spark' is already created from the previous question.

```
from pyspark.sql.functions import from_json, to_json, col
```

```
from pyspark.sql.types import StructType, StructField, StringType
```

--- SOLUTION ---

a) Create DataFrame with a JSON string column

```
json_data = [  
    (1, '{"name": "Alice", "city": "New York"}'),  
    (2, '{"name": "Bob", "city": "Los Angeles"}')  
]  
json_df = spark.createDataFrame(json_data, ["id", "json_str"])  
print("a) DataFrame with a JSON string column:")  
json_df.show(truncate=False)
```

b) Convert the JSON string into Struct type

First, define the schema that matches the JSON structure

```
json_schema = StructType([  
    StructField("name", StringType(), True),  
    StructField("city", StringType(), True)  
])  
df_with_struct = json_df.withColumn("parsed_struct", from_json(col("json_str"), json_schema))  
print("\nb) DataFrame with JSON converted to a StructType column:")  
df_with_struct.printSchema()  
df_with_struct.show(truncate=False)
```

c) Extract the Data from JSON and create them as new columns

```
df_extracted = df_with_struct.withColumn("name", col("parsed_struct.name")) \  
    .withColumn("city", col("parsed_struct.city"))  
print("\nc) DataFrame with JSON data extracted into new columns:")  
df_extracted.select("id", "name", "city").show()
```

d) Convert Struct type to JSON string

```
df_converted_back = df_extracted.withColumn("new_json_str", to_json(struct("name", "city")))  
print("\nd) DataFrame with columns converted back to a JSON string:")  
df_converted_back.select("id", "new_json_str").show(truncate=False)
```

- | | | |
|---|--|----|
| 1 | Create a data frame containing today's date, date 2022-01-31, date 2021-03-22, date 2024-01-31
Add 5 days to each date and display the result.
Display the new dates after subtracting 10 days from each date.
For each date, display year, month, dayofweek, dayofmonth, dayofyear, next_day, weekofyear | 20 |
| 2 | Refer to the employee.json file. Perform the following operations:
Print the names of employees above 25 years of age.
Print the number of employees of different ages. | 20 |
| 3 | Viva | 5 |
| 4 | Journal | 5 |

1.

Step 1: Install PySpark and set up the Spark Session

```
!pip install pyspark
```

```
from pyspark.sql import SparkSession
```

```
from pyspark.sql.functions import col, lit, to_date, date_add, date_sub, year, month, dayofweek, dayofmonth, dayofyear, next_day, weekofyear, current_date
```

```
spark = SparkSession.builder.appName("PracticalExam_Slip5_Q1").getOrCreate()
```

```
print("--- Spark Session Created ---\n")
```

--- SOLUTION ---

a) Create a DataFrame with dates

We create the DataFrame from string literals and convert them to DateType

```
date_data = ["2022-01-31", "2021-03-22", "2024-01-31"]
```

```
df_dates = spark.createDataFrame(date_data, "string").withColumnRenamed("value", "date_str")
```

```
df_dates = df_dates.union(spark.range(1).select(current_date().cast("string").alias("date_str"))) # Add today's date
```

```
df = df_dates.withColumn("date", to_date(col("date_str")))
```

```
print("a) Original DataFrame with dates:")
```

```
df.show()
```

b) Add 5 days to each date

```
print("\nb) Dates after adding 5 days:")
```

```
df_plus_5 = df.withColumn("date_plus_5", date_add(col("date"), 5))
```

```
df_plus_5.show()
```

c) Display the new dates after subtracting 10 days from each date

```
print("\nc) Dates after subtracting 10 days:")

df_minus_10 = df.withColumn("date_minus_10", date_sub(col("date"), 10))

df_minus_10.show()
```

d) Display various parts for each date

```
print("\nd) Various date parts:")

df_parts = df.withColumn("year", year(col("date"))) \
    .withColumn("month", month(col("date"))) \
    .withColumn("dayofweek", dayofweek(col("date"))) \
    .withColumn("dayofmonth", dayofmonth(col("date"))) \
    .withColumn("dayofyear", dayofyear(col("date"))) \
    .withColumn("next_day", next_day(col("date"), "Sunday")) \
    .withColumn("weekofyear", weekofyear(col("date")))

df_parts.show()
```

2.

Assuming the SparkSession 'spark' is already created from the previous question.

```
from pyspark.sql.functions import col
```

--- SETUP ---

Create the employee.json file in the Colab environment

```
json_content = """
{"name": "Michael", "age": 30}
{"name": "Andy", "age": 24}
{"name": "Justin", "age": 28}
{"name": "Berta", "age": 35}
{"name": "David", "age": 28}
"""
```

```
with open("employee.json", "w") as f:
```

```
    f.write(json_content)
```

```
print("--- employee.json file created successfully ---\n")
```

--- SOLUTION ---

Read the JSON file into a DataFrame

```
emp_df = spark.read.json("employee.json")
```

```
print("Original Employee DataFrame:")
```

```
emp_df.show()
```

```
# a) Print the names of employees above 25 years of age
```

```
print("\na) Names of employees older than 25:")
```

```
emp_df.filter(col("age") > 25).select("name").show()
```

```
# b) Print the number of employees of different ages
```

```
print("\nb) Number of employees for each age:")
```

```
emp_df.groupBy("age").count().show()
```

1	Create two dataframes one for employee and other for dept. Perform a) Left anti join b) Self join c) Left semi join	20
2	a) Create two case classes – Student and Address b) Create schema from these case classes	20
3	Viva	5
4	Journal	5

1.

Step 1: Install PySpark and set up the Spark Session

```
!pip install pyspark
```

```
from pyspark.sql import SparkSession
```

```
from pyspark.sql.functions import col
```

```
spark = SparkSession.builder.appName("PracticalExam_Slip6_Q1").getOrCreate()
```

```
print("--- Spark Session Created ---\n")
```

--- SETUP ---

Create the two DataFrames

```
emp_data = [(1, "Smith", 10), (2, "Rose", 20), (3, "Williams", 10), (4, "Jones", 30), (5, "Brown", 50)]
```

```
dept_data = [("Finance", 10), ("Marketing", 20), ("Sales", 30), ("IT", 40)]
```

```
emp_df = spark.createDataFrame(emp_data, ["emp_id", "name", "dept_id"])
```

```
dept_df = spark.createDataFrame(dept_data, ["dept_name", "dept_id"])
```

```
print("Employee DataFrame:")
```

```
emp_df.show()
```

```
print("Department DataFrame:")
```

```
dept_df.show()
```

--- SOLUTION ---

a) Left Anti Join

This join returns only the rows from the left DataFrame that do not have a match in the right DataFrame.

```
print("\na) Left Anti Join (Employees in departments not in the dept table):")
```

```
emp_df.join(dept_df, on="dept_id", how="left_anti").show()
```

b) Self Join

This is joining a DataFrame to itself. You must use aliases to distinguish them.

Example: Find employees who have the same department ID.

```
print("\nb) Self Join (Find pairs of employees in the same department):")
```

```
df1 = emp_df.alias("df1")
```

```
df2 = emp_df.alias("df2")
```

We add df1.emp_id < df2.emp_id to avoid duplicate pairs and self-joins.

```
self_join_df = df1.join(df2, on="dept_id") \
    .where(col("df1.emp_id") < col("df2.emp_id")) \
```

```
    .select(col("df1.name").alias("emp1"), col("df2.name").alias("emp2"), "dept_id")
```

```
self_join_df.show()
```

c) Left Semi Join

```
# This join is similar to an inner join, but it only returns the columns from the left DataFrame.
print("\nc) Left Semi Join (Employees in departments that exist in the dept table):")
emp_df.join(dept_df, on="dept_id", how="left_semi").show()
```

2.

```
# Assuming the SparkSession 'spark' is already created from the previous question.
```

```
from pyspark.sql import Row
```

```
from pyspark.sql.types import StructType, StructField, StringType, IntegerType
```

```
# --- EXPLANATION ---
```

```
# "Case classes" are a feature of the Scala language. They provide a concise way to define
```

```
# classes that are primarily used for holding data. Spark can automatically infer a schema from them in Scala.
```

```
#
```

```
# The direct equivalent in Python is to use a standard Python class or, more commonly,
```

```
# to define the schema explicitly using StructType and StructField.
```

```
# This practical demonstrates the PySpark way of achieving the same goal.
```

```
# --- SOLUTION ---
```

```
# a) Equivalent of creating "case classes"
```

```
# In Python, we can represent the data structure using standard classes or dictionaries.
```

```
# Here, we'll represent the data as nested Row objects, which is a common pattern.
```

```
student_data = [
```

```
    Row(name="John", age=20, address=Row(city="New York", zip_code="10001")),
```

```
    Row(name="Jane", age=22, address=Row(city="Los Angeles", zip_code="90001"))
```

```
]
```

```
# b) Create schema from the data structure
```

```
# We explicitly define the schema to match our data structure.
```

```
address_schema = StructType([
```

```
    StructField("city", StringType(), True),
```

```
    StructField("zip_code", StringType(), True)
```

```
])
```

```
student_schema = StructType([
```

```
    StructField("name", StringType(), True),
```

```
    StructField("age", IntegerType(), True),
```

```
    StructField("address", address_schema, True)
```

```
])
```

```
print("b) Explicitly created schema:")
```

```
student_schema.prettyJson()
```

```
# Now, create a DataFrame using this data and schema
```

```
student_df = spark.createDataFrame(student_data, student_schema)
```

```
print("\nDataFrame created from the schema:")
```

```
student_df.show(truncate=False)
```

```
student_df.printSchema()
```

1	Create a data frame with data that follows the below given schema emp_id, dept, properties (a structure containing salary and location) Return the map keys from spark SQL for this data frame	20
2	For the following employee data showing name, dept and salary, perform the given operations: Data: ("James", "Sales", 3000), ("Michael", "Sales", 4600), ("Robert", "Sales", 4100), ("Maria", "Finance", 3000), ("James", "Sales", 3000), ("Scott", "Finance", 3300), ("Jen", "Finance", 3900), ("Jeff", "Marketing", 3000), ("Kumar", "Marketing", 2000), ("Saif", "Sales", 4100), ("Jason", "Sales", 9000), ("Alice", "Finance", 3700), ("Jenniffer", "Finance", 8900), ("Jenson", "Marketing", 9000) a) Create a data frame for the above data b) Find the highest salary value c) Find the lowest salary value d) Find the standard deviation for the salary	20
3	Viva	5
4	Journal	5

1.

Step 1: Install PySpark and set up the Spark Session

!pip install pyspark

from pyspark.sql import SparkSession, Row

from pyspark.sql.functions import col

spark = SparkSession.builder.appName("PracticalExam_Slip7_Q1").getOrCreate()

print("--- Spark Session Created ---\n")

--- EXPLANATION ---

The question asks to "Return the map keys", but the `properties` column is a struct, not a map.

A struct has fixed fields (like keys), while a map can have arbitrary key-value pairs.

The correct interpretation is to access the fields of the struct, not "map keys".

--- SETUP ---

Create a DataFrame with a struct column

data = [

 (1, "Finance", Row(salary=90000, location="New York")),

 (2, "Marketing", Row(salary=80000, location="Chicago")),

 (3, "Sales", Row(salary=120000, location="New York"))

]

df = spark.createDataFrame(data, ["emp_id", "dept", "properties"])

print("Original DataFrame with Struct column:")

df.show(truncate=False)

df.printSchema()

```
# --- SOLUTION ---
# Access the fields of the 'properties' struct
print("\nAccessing the fields ('keys') of the 'properties' struct:")
df.select(
    col("emp_id"),
    col("properties.salary"),
    col("properties.location")
).show()
```

2.

```
# Assuming the SparkSession 'spark' is already created from the previous question.
from pyspark.sql.functions import max, min, stddev
```

```
# --- SETUP ---
```

```
# The provided data has a typo `Jason`, which has been corrected to `("Jason")`.
```

```
employee_data = [
    ("James", "Sales", 3000), ("Michael", "Sales", 4600),
    ("Robert", "Sales", 4100), ("Maria", "Finance", 3000),
    ("James", "Sales", 3000), ("Scott", "Finance", 3300),
    ("Jen", "Finance", 3900), ("Jeff", "Marketing", 3000),
    ("Kumar", "Marketing", 2000), ("Saif", "Sales", 4100),
    ("Jason", "Sales", 9000), ("Alice", "Finance", 3700),
    ("Jenniffer", "Finance", 8900), ("Jenson", "Marketing", 9000)
]
columns = ["name", "department", "salary"]
```

```
# --- SOLUTION ---
```

```
# a) Create a data frame for the above data
```

```
emp_df = spark.createDataFrame(employee_data, columns)
print("\na) Employee DataFrame:")
emp_df.show()
```

```
# b) Find the highest salary value
```

```
print("\nb) Highest Salary:")
emp_df.select(max("salary").alias("highest_salary")).show()
```

```
# c) Find the lowest salary value
```

```
print("\nc) Lowest Salary:")
emp_df.select(min("salary").alias("lowest_salary")).show()
```

```
# d) Find the standard deviation for the salary
```

```
print("\nd) Standard Deviation of Salary:")
emp_df.select(stddev("salary").alias("stddev_salary")).show()
```


1	Create a data frame with a nested array column. Perform the following operations: a) Flatten nested array b) Explode nested array c) Convert array of string to string column.	20
2	a) Create data frame with a column that contains JSON string. b) Convert the JSON string into Struct type or Map type. c) Extract the Data from JSON and create them as new columns. d) Convert MapType or Struct type to JSON string	20
3	Viva	5
4	Journal	5

1.

Step 1: Install PySpark and set up the Spark Session

```
!pip install pyspark
```

```
from pyspark.sql import SparkSession
```

```
from pyspark.sql.functions import col, flatten, explode, concat_ws
```

```
from pyspark.sql.types import StructType, StructField, StringType, ArrayType
```

```
spark = SparkSession.builder.appName("PracticalExam_Slip8_Q1").getOrCreate()
```

```
print("--- Spark Session Created ---\n")
```

--- SETUP ---

Create the DataFrame with a nested array

```
data = [
```

```
    ("James", [ ["Java", "Scala", "C++"], ["Spark", "Java"] ]),
```

```
    ("Michael", [ ["Spark", "Java", "C++"], ["Spark", "Java"] ]),
```

```
    ("Robert", [ ["CSharp", "VB"], ["Spark", "Python"] ])
```

```
]
```

```
schema = StructType([
```

```
    StructField("name", StringType(), True),
```

```
    StructField("subjects", ArrayType(ArrayType(StringType())), True)
```

```
])
```

```
df = spark.createDataFrame(data, schema)
```

```
print("Original DataFrame with nested array:")
```

```
df.show(truncate=False)
```

--- SOLUTION ---

a) Flatten nested array

```
print("\na) Flattened nested array:")
```

```
df_a = df.withColumn("subjects_flat", flatten(col("subjects")))
```

```
df_a.show(truncate=False)
```

b) Explode nested array

```
print("\nb) Exploded nested array to rows:")
```

To get individual elements, you must flatten first.

```
df_b = df_a.withColumn("subject", explode(col("subjects_flat")))
```

```
df_b.show(truncate=False)
```

c) Convert array of string to string column

```
print("\nc) Converted array to a single string column:")
```

```
df_c = df_a.withColumn("subjects_string", concat_ws(" ", col("subjects_flat")))
```

```
df_c.show(truncate=False)
```

2.

Assuming the SparkSession 'spark' is already created from the previous question.

```
from pyspark.sql.functions import from_json, to_json, col, struct
```

```
from pyspark.sql.types import StructType, StructField, StringType
```

--- SOLUTION ---

a) Create DataFrame with a JSON string column

```
json_data = [
```

```
    (1, '{"name": "Alice", "city": "New York"}'),
```

```
    (2, '{"name": "Bob", "city": "Los Angeles"}')
```

```
]
```

```
json_df = spark.createDataFrame(json_data, ["id", "json_str"])
```

```
print("a) DataFrame with a JSON string column:")
```

```
json_df.show(truncate=False)
```

b) Convert the JSON string into Struct type

```
# Define the schema that matches the JSON structure
```

```
json_schema = StructType([
    StructField("name", StringType(), True),
    StructField("city", StringType(), True)
])

df_with_struct = json_df.withColumn("parsed_struct", from_json(col("json_str"), json_schema))

print("\nb) DataFrame with JSON converted to a StructType column:")

df_with_struct.printSchema()

df_with_struct.show(truncate=False)
```

```
# c) Extract the Data from JSON and create them as new columns
```

```
df_extracted = df_with_struct.withColumn("name", col("parsed_struct.name")) \
    .withColumn("city", col("parsed_struct.city"))

print("\nc) DataFrame with JSON data extracted into new columns:")

df_extracted.select("id", "name", "city").show()
```

```
# d) Convert Struct type to JSON string
```

```
df_converted_back = df_extracted.withColumn("new_json_str", to_json(struct("name", "city")))

print("\nd) DataFrame with columns converted back to a JSON string:")

df_converted_back.select("id", "new_json_str").show(truncate=False)
```

1	Create a Spark RDD using 5 different Functions	20
2	Write example for following Spark RDD Actions: a. aggregate b. treeAggregate c. fold d. reduce e. collect	20
3	Viva	5
4	Journal	5

1.

Step 1: Install PySpark and set up the Spark Session

!pip install pyspark

from pyspark.sql import SparkSession

spark = SparkSession.builder.appName("PracticalExam_Slip9_Q1").getOrCreate()

sc = spark.sparkContext

print("--- Spark Session Created ---\n")

--- SETUP ---

Create a sample text file

with open("sample_rdd.txt", "w") as f:

f.write("line one\n")

f.write("line two\n")

f.write("line three\n")

print("--- Sample file for RDD creation created ---\n")

--- SOLUTION ---

print("--- Creating RDDs using 5 different methods ---\n")

1. Using sc.parallelize() on a list

list_data = [1, 2, 3, 4, 5]

rdd1 = sc.parallelize(list_data)

print("1. RDD from a Python list (parallelize):")

print(rdd1.collect())

2. Using sc.textFile() to read a text file

```
rdd2 = sc.textFile("sample_rdd.txt")
print("\n2. RDD from a text file (textFile):")
print(rdd2.collect())
```

```
# 3. Using sc.range()
rdd3 = sc.range(1, 6) # Creates an RDD with elements 1, 2, 3, 4, 5
print("\n3. RDD from a range (range):")
print(rdd3.collect())
```

```
# 4. By transforming an existing RDD (e.g., using map)
rdd4 = rdd1.map(lambda x: x * x)
print("\n4. RDD by transforming another RDD (map):")
print(rdd4.collect())
```

```
# 5. From a DataFrame
df = spark.createDataFrame([("a", 1), ("b", 2)], ["letter", "number"])
rdd5 = df.rdd
print("\n5. RDD from a DataFrame (.rdd):")
print(rdd5.collect())
```

2.

Assuming the SparkContext 'sc' is already created from the previous question.

--- SETUP ---

```
rdd = sc.parallelize([1, 2, 3, 4, 5])
print("Using the following RDD for actions:", rdd.collect())
```

--- SOLUTION ---

a) aggregate

Action: Sums elements and adds an initial value to each partition and then to the final result.

(sum_of_elements + initial_value * num_partitions) + initial_value

Here: (1+2+3+4+5) -> 15. Let's assume 2 partitions.

Partition 1: 1+2 -> 3. Partition 2: 3+4+5 -> 12.

Add initial value 10 to each partition sum: $(3+10) + (12+10) = 45$

```
seqOp = (lambda x, y: x + y)
```

```
combOp = (lambda x, y: x + y)
```

```
agg_result = rdd.aggregate(0, seqOp, combOp) # Using 0 as initial value is same as reduce
```

```
print("\na) aggregate (sum):", agg_result)
```

b) treeAggregate

Similar to aggregate but performs aggregation in a tree-like pattern, which is more efficient for large datasets.

```
tree_agg_result = rdd.treeAggregate(0, seqOp, combOp)
```

```
print("\nb) treeAggregate (sum):", tree_agg_result)
```

c) fold

Similar to reduce but takes a "zero value" to be used for the initial call in each partition.

```
fold_result = rdd.fold(0, lambda x, y: x + y)
```

```
print("\nc) fold (sum):", fold_result)
```

d) reduce

Aggregates the elements of the RDD using a specified commutative and associative binary operator.

```
reduce_result = rdd.reduce(lambda x, y: x + y)
```

```
print("\nd) reduce (sum):", reduce_result)
```

e) collect

Returns all the elements of the RDD as a list to the driver program.

```
collect_result = rdd.collect()
```

```
print("\ne) collect:", collect_result)
```

1	Write example for following Spark RDD Actions: a. count b. countApproxDistinct c. first d. top e. Min	20
2	Write Spark Pair RDD Functions.	20
3	Viva	5
4	Journal	5

1.

Step 1: Install PySpark and set up the Spark Session

```
!pip install pyspark
```

```
from pyspark.sql import SparkSession
```

```
spark = SparkSession.builder.appName("PracticalExam_Slip10_Q1").getOrCreate()
```

```
sc = spark.sparkContext
```

```
print("--- Spark Session Created ---\n")
```

```
# --- SETUP ---
```

```
rdd = sc.parallelize([2, 5, 1, 3, 4, 2, 5])
```

```
print("Using the following RDD for actions:", rdd.collect())
```

```
# --- SOLUTION ---
```

```
# a) count
```

```
# Returns the number of elements in the RDD.
```

```
count_result = rdd.count()
```

```
print("\na) count:", count_result)
```

```
# b) countApproxDistinct
```

```
# Returns the approximate number of distinct elements. Useful for large datasets.
```

```
approx_distinct_result = rdd.countApproxDistinct()
```

```
print("\nb) countApproxDistinct:", approx_distinct_result)
```

```
# Note: The exact distinct count is 5 (1, 2, 3, 4, 5)
```

```
# c) first
```

```
# Returns the first element of the RDD.
```

```
first_result = rdd.first()
```

```
print("\nc) first:", first_result)
```

```
# d) top
```

```
# Returns the top n elements from an RDD, ordered in descending order.
```

```
top_3_result = rdd.top(3)
```

```
print("\nd) top(3):", top_3_result)
```

```
# e) min
```

```
# Returns the minimum element of the RDD.
```

```
min_result = rdd.min()
```

```
print("\ne) min:", min_result)
```

2.

Assuming the SparkContext 'sc' is already created from the previous question.

--- SETUP ---

A Pair RDD is an RDD where each element is a key-value tuple.

Let's create a sample Pair RDD.

```
data = [("apple", 1), ("banana", 2), ("apple", 3), ("orange", 4), ("banana", 5)]
```

```
pair_rdd = sc.parallelize(data)
```

```
print("Using the following Pair RDD:", pair_rdd.collect())
```

--- SOLUTION ---

Here are examples of common Pair RDD functions.

1. reduceByKey()

Merges the values for each key using an associative and commutative reduce function.

```
print("\n1. reduceByKey (sum of values for each key):")
```

```
reduced_rdd = pair_rdd.reduceByKey(lambda a, b: a + b)
```

```
print(reduced_rdd.collect())
```

2. groupByKey()

Groups the values for each key in the RDD into a single sequence.

```
print("\n2. groupByKey:")
```

```
grouped_rdd = pair_rdd.groupByKey()
```

The result contains an iterable object, so we map it to a list for printing.

```
print(grouped_rdd.mapValues(list).collect())
```

3. sortByKey()

Sorts the RDD by key.

```
print("\n3. sortByKey (ascending):")
```

```
sorted_rdd = pair_rdd.sortByKey()
```

```
print(sorted_rdd.collect())
```

4. keys() and values()

Return an RDD of just the keys or just the values.

```
print("\n4. keys() and values():")
```

```
keys_rdd = pair_rdd.keys()
```

```
values_rdd = pair_rdd.values()
```

```
print("Keys:", keys_rdd.collect())
```

```
print("Values:", values_rdd.collect())
```

5. join()

Joins two Pair RDDs based on their keys.

```
other_data = [("apple", "red"), ("orange", "orange"), ("grape", "purple")]
```

```
other_pair_rdd = sc.parallelize(other_data)
```

```
print("\n5. join:")
```

```
joined_rdd = pair_rdd.join(other_pair_rdd)
```

```
print(joined_rdd.collect())
```


1	Get new dates by adding 4 days, and subtracting 7 days in below dates "2020-01-02","2023-01-15","2025-01-30"	20
2	Use the Operation Read CSV file on RDD with Scala operation	20
3	Viva	5
4	Journal	5

1.

Step 1: Install PySpark and set up the Spark Session

```
!pip install pyspark
```

```
from pyspark.sql import SparkSession
```

```
from pyspark.sql.functions import col, to_date, date_add, date_sub
```

```
spark = SparkSession.builder.appName("PracticalExam_Slip11_Q1").getOrCreate()
```

```
print("--- Spark Session Created ---\n")
```

```
# --- SETUP ---
```

```
date_data = ["2020-01-02", "2023-01-15", "2025-01-30"]
```

```
df = spark.createDataFrame(date_data, "string").withColumnRenamed("value", "date_str")
```

```
# Convert strings to DateType
```

```
df = df.withColumn("date", to_date(col("date_str")))
```

```
print("Original DataFrame with dates:")
```

```
df.show()
```

```
# --- SOLUTION ---
```

```
# Add 4 days to each date
```

```
df_plus_4 = df.withColumn("date_plus_4_days", date_add(col("date"), 4))
```

```
print("\nDates after adding 4 days:")
```

```
df_plus_4.show()
```

```
# Subtract 7 days from each date
```

```
df_minus_7 = df.withColumn("date_minus_7_days", date_sub(col("date"), 7))
```

```
print("\nDates after subtracting 7 days:")
```

```
df_minus_7.show()
```

2.

Assuming the SparkContext 'sc' and SparkSession 'spark' are already created.

```
# --- EXPLANATION ---
```

```
# The question asks to use a "Scala operation". In PySpark, we use Python operations (like lambda functions).
```

```
# The most common way to read a CSV is with a DataFrame, but to fulfill the "on RDD" requirement,
```

```
# we will read the file as a text RDD and then parse it.
```

```
# --- SETUP ---
```

```
# Create a sample CSV file
```

```
with open("student_data.csv", "w") as f:
```

```
    f.write("id,name,score\n")
```

```
    f.write("1,Alice,85\n")
```

```
    f.write("2,Bob,90\n")
```

```
    f.write("3,Cathy,78\n")
```

```
print("--- student_data.csv created successfully ---\n")
```

--- SOLUTION ---

1. Read the CSV file into a text RDD

```
text_rdd = sc.textFile("student_data.csv")
print("RDD as raw text lines:")
print(text_rdd.collect())
```

2. Get the header and filter it out

```
header = text_rdd.first()
data_rdd = text_rdd.filter(lambda line: line != header)
print("\nRDD after removing the header:")
print(data_rdd.collect())
```

3. Use an RDD operation (map) to parse the data

This is the equivalent of a "Scala operation" in PySpark.

```
parsed_rdd = data_rdd.map(lambda line: line.split(","))
# Convert score to an integer
parsed_rdd = parsed_rdd.map(lambda parts: (int(parts[0]), parts[1], int(parts[2])))
```

```
print("\nRDD after parsing and transforming with a map operation:")
print(parsed_rdd.collect())
```

Example: Filter for scores above 80

```
high_scores_rdd = parsed_rdd.filter(lambda x: x[2] > 80)
print("\nResult of another operation (filter for score > 80):")
print(high_scores_rdd.collect())
```

1	Create table as follows containing array and map operations	20
	<pre> +-----+-----+-----+ name knownLanguages properties +-----+-----+-----+ James [Java, Scala] {hair -> black, eye -> brown} Michael [Spark, Java, null] {hair -> brown, eye -> null} Robert [CSharp,] {hair -> red, eye -> } Washington null null Jefferson [] {} +-----+-----+-----+ </pre>	
2	Find current timestamp and hour, Minute, second separately for today's date	20
3	Viva	5
4	Journal	5

```

1.
# Step 1: Install PySpark and set up the Spark Session
!pip install pyspark
from pyspark.sql import SparkSession
from pyspark.sql.types import StructType, StructField, StringType, ArrayType, MapType

spark = SparkSession.builder.appName("PracticalExam_Slip12_Q1").getOrCreate()
print("--- Spark Session Created ---\n")

# --- EXPLANATION ---
# The table in the question is poorly formatted but describes a DataFrame with three columns:
# 'name' (String), 'knownLanguages' (Array of Strings), and 'properties' (Map of String to String).
# We will create the data based on this interpretation.

# --- SOLUTION ---

# Define the data, handling nulls and empty values from the question
data = [
    ("James", ["Java", "Scala"], {"hair": "black", "eye": "brown"}),
    ("Michael", ["Spark", "Java", None], {"hair": "brown", "eye": None}),
    ("Robert", ["CSharp", None], {"hair": "red", "eye": None})
]

# Define the schema for the DataFrame
schema = StructType([
    StructField("name", StringType(), True),
    StructField("knownLanguages", ArrayType(StringType()), True),
    StructField("properties", MapType(StringType(), StringType()), True)
])

# Create the DataFrame
df = spark.createDataFrame(data, schema)

print("Created DataFrame:")
df.show(truncate=False)
df.printSchema()

```

2.

Assuming the SparkSession 'spark' is already created from the previous question.

```
from pyspark.sql.functions import current_timestamp, hour, minute, second, col
```

--- SOLUTION ---

Create a dummy DataFrame with one row to demonstrate the functions

```
df = spark.range(1)
```

Add a column with the current timestamp

```
df_with_ts = df.withColumn("current_timestamp", current_timestamp())
```

```
print("DataFrame with current timestamp:")
```

```
df_with_ts.show(truncate=False)
```

Extract hour, minute, and second into separate columns

```
time_parts_df = df_with_ts.withColumn("hour", hour(col("current_timestamp"))) \
```

```
    .withColumn("minute", minute(col("current_timestamp"))) \
```

```
    .withColumn("second", second(col("current_timestamp")))
```

```
print("\nTimestamp with hour, minute, and second extracted:")
```

```
time_parts_df.select("current_timestamp", "hour", "minute", "second").show(truncate=False)
```

1	Write a Maven dependencies for writing and Reading Avro Data File	20
2	Create the following two data frames and apply Inner and Right Outer <pre> +-----+-----+-----+-----+-----+-----+ emp_id name superior_emp_id year_joined emp_dept_id gender +-----+-----+-----+-----+-----+-----+ 1 Smith -1 2018 10 M 2 Rose 1 2010 20 M 3 Williams 1 2010 10 M 4 Jones 2 2005 10 F 5 Brown 2 2010 40 6 Brown 2 2010 50 +-----+-----+-----+-----+-----+-----+ </pre> join. 20M <pre> +-----+-----+ dept_name dept_id +-----+-----+ Finance 10 Marketing 20 Sales 30 IT 40 +-----+-----+ </pre>	20
3	Viva	5
4	Journal	5

1.

1. To launch pyspark from the command line

Bash:

```
pyspark --packages org.apache.spark:spark-avro_2.12:3.5.1
```

2. To configure it within a Colab notebook or Python script:

```
from pyspark.sql import SparkSession
```

```

spark = SparkSession.builder \
    .appName("AvroExample") \
    .config("spark.jars.packages", "org.apache.spark:spark-avro_2.12:3.5.1") \
    .getOrCreate()

```

```
print("SparkSession created with Avro package.")
```

2.

Step 1: Install PySpark and set up the Spark Session

```
!pip install pyspark
```

```
from pyspark.sql import SparkSession
```

```
from pyspark.sql.functions import col
```

```
spark = SparkSession.builder.appName("PracticalExam_Slip13_Q2").getOrCreate()
```

```
print("--- Spark Session Created ---\n")
```

```
# --- SETUP ---
```

```
# Create the two DataFrames based on the provided tables.
```

```
emp_data = [
```

```
    (1, "Smith", None, 2018, 10, "M"),
```

```

(2, "Rose", 1, 2010, 20, "M"),
(3, "Williams", 1, 2010, 10, "M"),
(4, "Jones", 2, 2005, 10, "F"),
(5, "Brown", 2, 2010, 40, None),
(6, "Brown", 2, 2010, 50, None) # Corrected from the original slip's typo
]

emp_columns = ["emp_id", "name", "superior_emp_id", "year_joined", "emp_dept_id", "gender"]
emp_df = spark.createDataFrame(emp_data, emp_columns)

dept_data = [("Finance", 10), ("Marketing", 20), ("Sales", 30), ("IT", 40)]
dept_columns = ["dept_name", "dept_id"]
dept_df = spark.createDataFrame(dept_data, dept_columns)

print("Employee DataFrame:")
emp_df.show()
print("Department DataFrame:")
dept_df.show()

# --- SOLUTION ---
# We need to join on the department ID. The columns have different names,
# so we need to specify the join condition explicitly.
join_condition = emp_df.emp_dept_id == dept_df.dept_id

# a) Inner Join
print("\na) Inner Join Result:")
emp_df.join(dept_df, join_condition, "inner").show()

# b) Right Outer Join
print("\nb) Right Outer Join Result:")
emp_df.join(dept_df, join_condition, "right_outer").show()

```